

3.1 The Backpropagation Algorithm (Rumelhart et al., 1986) ^{23.}

Formulation for a feed-forward 2-layer network of sigmoid units, the
stochastic version

Idea: Gradient descent over the entire vector of network weights.

Initialize all weights to small random numbers.

Until satisfied, // *stopping criterion* to be (later) defined
for each training example,

1. input the training example to the network, and compute the network outputs
2. for each output unit k :
$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$
3. for each hidden unit h :
$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{kh} \delta_k$$
4. update each network weight w_{ji} :
$$w_{ji} \leftarrow w_{ji} + \Delta w_{ji} \text{ where } \Delta w_{ji} = \eta \delta_j x_{ji},$$

and x_{ji} is the i th input to unit j .

Derivation of the Backpropagation rule,

(following [Tom Mitchell, 1997], pag. 101–103)

Notations:

x_{ji} : the i th input to unit j ;
(j could be either hidden or output unit)

w_{ji} : the weight associated with the i th input to unit j

$$net_j = \sum_i w_{ji} x_{ji}$$

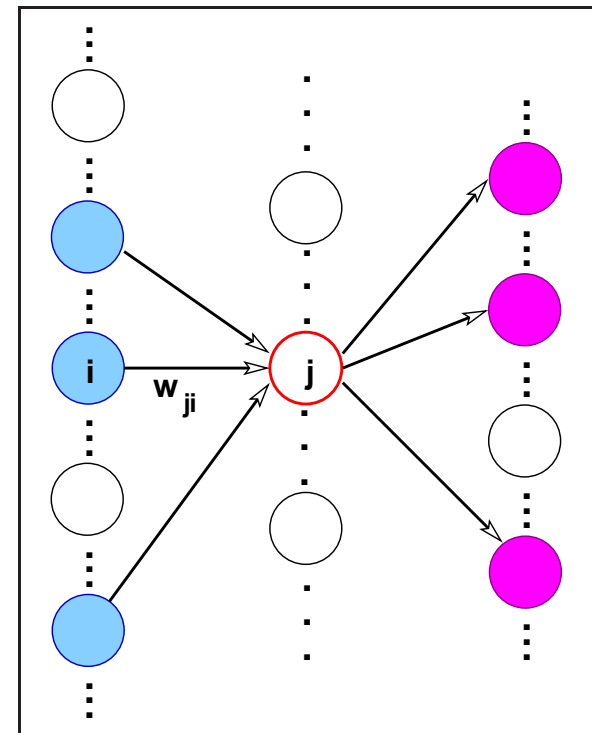
σ : the sigmoid function

o_j : the output computed by unit j ; ($o_j = \sigma(net_j)$)

outputs: the set of units in the final layer of the network

Downstream(j): the set of units whose immediate inputs include the output of unit j

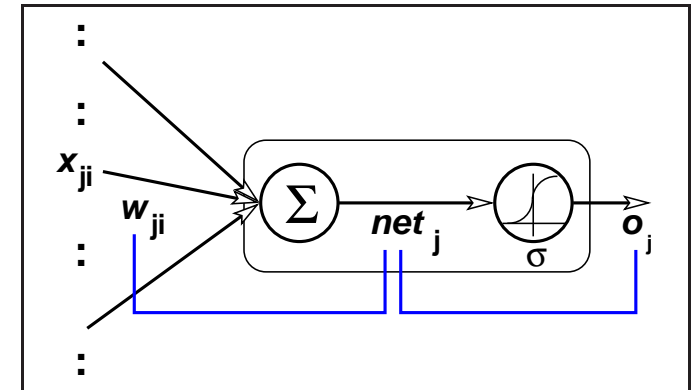
E_d : the training error on the example d (summing over all of the network output units)



Legend: in **magenta color**, units belonging to *Downstream*(j)

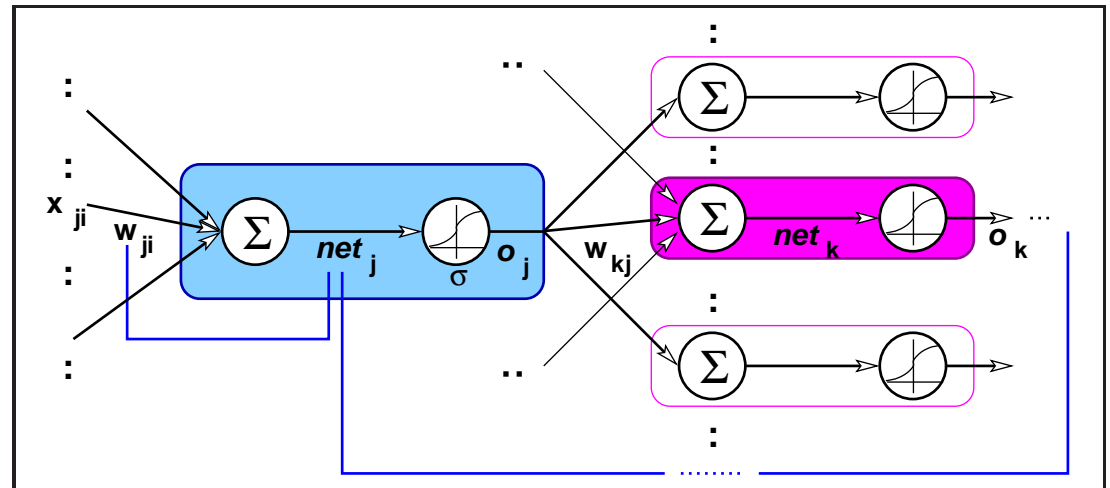
Preliminaries

$$E_d(\vec{w}) = \frac{1}{2} \sum_{k \in \text{outputs}} (t_k - o_k)^2 = \frac{1}{2} \sum_{k \in \text{outputs}} (t_k - \sigma(\text{net}_k))^2$$



Common stuff for both hidden and output units:

$$\begin{aligned} \text{net}_j &= \sum_i w_{ji} x_{ji} \\ \Rightarrow \frac{\partial E_d}{\partial w_{ji}} &= \frac{\partial E_d}{\partial \text{net}_j} \frac{\partial \text{net}_j}{\partial w_{ji}} = \frac{\partial E_d}{\partial \text{net}_j} x_{ji} \\ \Rightarrow \Delta w_{ji} &\stackrel{\text{def}}{=} -\eta \frac{\partial E_d}{\partial w_{ji}} = -\eta \frac{\partial E_d}{\partial \text{net}_j} x_{ji} \end{aligned}$$



Note: In the sequel we will use the notation: $\delta_j = -\frac{\partial E_d}{\partial \text{net}_j} \Rightarrow \Delta w_{ji} = \eta \delta_j x_{ji}$

Stage/Case 1: Computing the increments (Δ) for output unit weights

$$\frac{\partial E_d}{\partial net_j} = \frac{\partial E_d}{\partial o_j} \frac{\partial o_j}{\partial net_j}$$


$$\frac{\partial o_j}{\partial net_j} = \frac{\partial \sigma(net_j)}{\partial net_j} = o_j(1 - o_j)$$

$$\begin{aligned} \frac{\partial E_d}{\partial o_j} &= \frac{\partial}{\partial o_j} \frac{1}{2} \sum_{k \in outputs} (t_k - o_k)^2 \\ &= \frac{\partial}{\partial o_j} \frac{1}{2} (t_j - o_j)^2 = \frac{1}{2} 2(t_j - o_j) \frac{\partial (t_j - o_j)}{\partial o_j} \\ &= -(t_j - o_j) \end{aligned}$$

$$\Rightarrow \frac{\partial E_d}{\partial net_j} = -(t_j - o_j) o_j (1 - o_j) = -o_j (1 - o_j) (t_j - o_j)$$

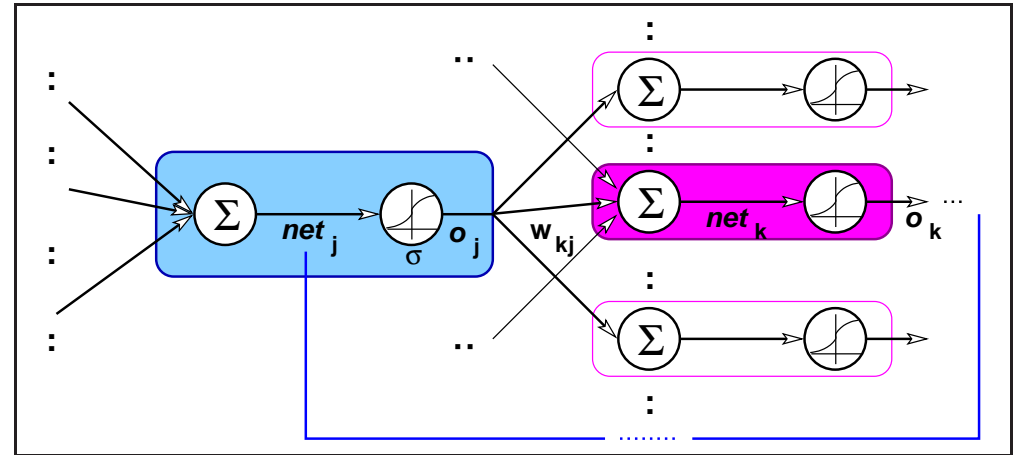
$$\Rightarrow \delta_j \stackrel{not.}{=} -\frac{\partial E_d}{\partial net_j} = o_j (1 - o_j) (t_j - o_j)$$

$$\Rightarrow \Delta w_{ji} = \eta \delta_j x_{ji} = \eta o_j (1 - o_j) (t_j - o_j) x_{ji}$$



Stage/Case 2: Computing the increments (Δ) for hidden unit weights

$$\begin{aligned}
 \frac{\partial E_d}{\partial net_j} &= \sum_{k \in \text{Downstream}(j)} \frac{\partial E_d}{\partial net_k} \frac{\partial net_k}{\partial net_j} \\
 &= \sum_{k \in \text{Downstream}(j)} -\delta_k \frac{\partial net_k}{\partial net_j} \\
 &= \sum_{k \in \text{Downstream}(j)} -\delta_k \frac{\partial net_k}{\partial o_j} \frac{\partial o_j}{\partial net_j}
 \end{aligned}$$



$$= \sum_{k \in \text{Downstream}(j)} -\delta_k w_{kj} \frac{\partial o_j}{\partial net_j} = \sum_{k \in \text{Downstream}(j)} -\delta_k w_{kj} o_j (1 - o_j)$$

Therefore:

$$\begin{aligned}
 \delta_j &\stackrel{\text{not}}{=} -\frac{\partial E_d}{\partial net_j} = o_j (1 - o_j) \sum_{k \in \text{Downstream}(j)} \delta_k w_{kj} \\
 \Delta w_{ji} &\stackrel{\text{def}}{=} -\eta \frac{\partial E_d}{\partial w_{ji}} = -\eta \frac{\partial E_d}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}} = -\eta \frac{\partial E_d}{\partial net_j} x_{ji} = \eta \delta_j x_{ji} = \eta \left[o_j (1 - o_j) \sum_{k \in \text{Downstream}(j)} \delta_k w_{kj} \right] x_{ji}
 \end{aligned}$$


Convergence of Backpropagation for NNs of Sigmoid units

Nature of convergence

- The weights are initialized near zero; therefore, initial decision surfaces are near-linear.

Explanation: o_j is of the form $\sigma(\vec{w} \cdot \vec{x})$, therefore $w_{ji} \approx 0$ for all i, j ;

note that the graph of σ is approximately linear in the vicinity of 0.



- Increasingly non-linear functions are possible as training progresses
- Will find a **local**, not necessarily global error **minimum**.
In **practice**, often works well (can run multiple times).

More on Backpropagation

- Easily generalized to arbitrary directed graphs
- Training can take thousands of iterations → slow!
- Often include weight momentum α

$$\Delta w_{i,j}(n) = \eta \delta_j x_{ij} + \alpha \Delta w_{ij}(n-1)$$

Effect:

- speed up convergence (increase the step size in regions where the gradient is unchanging);
 - “keep the ball rolling” through local minima (or along flat regions) in the error surface
- Using network after training is very fast
 - Minimizes error over *training* examples;
Will it generalize well to subsequent examples?